

Derivation of the Adam Optimization Algorithm

MHJ

Notes 2018

Overview

Why Combine Momentum and RMSProp?

Motivation for Adam: Adaptive Moment Estimation (Adam) was introduced by Kingma & Ba (2014) to combine the benefits of momentum and RMSProp

- **Momentum:** Fast convergence by smoothing gradients (accelerates in long-term gradient direction).
- **Adaptive rates (RMSProp):** Per-dimension learning rate scaling for stability (handles different feature scales, sparse gradients).
- Adam uses both: maintains moving averages of both first moment (gradients) and second moment (squared gradients)
- Additionally, includes a mechanism to correct the bias in these moving averages (crucial in early iterations)
- Result: Adam is robust, achieves faster convergence with less tuning, and often outperforms SGD (with momentum) in practice

Adam: Exponential Moving Averages (Moments)

Adam maintains two moving averages at each time step t for each parameter w :

First moment (mean) m_t

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla L(w_t), \quad (\text{Momentum term})$$

Second moment (uncentered variance) v_t

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla L(w_t))^2, \quad (\text{RMS term})$$

with typical $\beta_1 = 0.9$, $\beta_2 = 0.999$. Initialize $m_0 = 0$, $v_0 = 0$.

These are *biased* estimators of the true first and second moment of the gradients, especially at the start (since m_0, v_0 are zero)

Adam: Bias Correction

To counteract initialization bias in m_t, v_t , Adam computes bias-corrected estimates

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}.$$

- When t is small, $1 - \beta_i^t \approx 0$, so \hat{m}_t, \hat{v}_t significantly larger than raw m_t, v_t , compensating for the initial zero bias.
- As t increases, $1 - \beta_i^t \rightarrow 1$, and \hat{m}_t, \hat{v}_t converge to m_t, v_t .
- Bias correction is important for Adam's stability in early iterations

Adam: Update Rule Derivation

Finally, Adam updates parameters using the bias-corrected moments:

$$w_{t+1} = w_t - \frac{\alpha}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t,$$

where ϵ is a small constant (e.g. 10^{-8}) to prevent division by zero

Breaking it down:

- 1 Compute gradient $\nabla L(w_t)$.
- 2 Update first moment m_t and second moment v_t (exponential moving averages).
- 3 Bias-correct: $\hat{m}_t = m_t / (1 - \beta_1^t)$, $\hat{v}_t = v_t / (1 - \beta_2^t)$.
- 4 Compute step: $\Delta w_t = \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$.
- 5 Update parameters: $w_{t+1} = w_t - \alpha \Delta w_t$.

This is the Adam update rule as given in the original paper

Adam vs. AdaGrad and RMSProp

- **AdaGrad:** Uses per-coordinate scaling like Adam, but no momentum. Tends to slow down too much due to cumulative history (no forgetting)
- **RMSProp:** Uses moving average of squared gradients (like Adam's v_t) to maintain adaptive learning rates, but does not include momentum or bias-correction.
- **Adam:** Effectively RMSProp + Momentum + Bias-correction
 - Momentum (m_t) provides acceleration and smoother convergence.
 - Adaptive v_t scaling moderates the step size per dimension.
 - Bias correction (absent in AdaGrad/RMSProp) ensures robust estimates early on.
- In practice, Adam often yields faster convergence and better tuning stability than RMSProp or AdaGrad alone

Adaptivity Across Dimensions

- Adam adapts the step size *per coordinate*: parameters with larger gradient variance get smaller effective steps, those with smaller or sparse gradients get larger steps.
- This per-dimension adaptivity is inherited from AdaGrad/RMSProp and helps handle ill-conditioned or sparse problems.
- Meanwhile, momentum (first moment) allows Adam to continue making progress even if gradients become small or noisy, by leveraging accumulated direction.
- **Example:** In a deep network, some weights may receive very noisy or infrequent gradients – Adam will keep their learning rate high (unlike AdaGrad which would have decayed it) and also smooth out the noise with momentum.