

# Derivations

RNNs and AEs

November 16, 2025

# Vanilla RNN: forward equations

- Hidden state update:

$$h_t = \phi(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

# Vanilla RNN: forward equations

- Hidden state update:

$$h_t = \phi(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

- Output:

$$y_t = W_{hy}h_t + b_y$$

# Vanilla RNN: forward equations

- Hidden state update:


$$h_t = \phi(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

- Output:

$$y_t = W_{hy}h_t + b_y$$

- Where  $\phi$  is e.g. tanh or ReLU.

# Unroll in time (visual)



# Backpropagation Through Time (BPTT) — derivation

Loss over sequence

$$\mathcal{L} = \sum_{t=1}^T \ell(y_t, \hat{y}_t)$$

$$\frac{\partial \mathcal{L}}{\partial W_{hh}} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial h_t} \frac{\partial h_t}{\partial W_{hh}}.$$

$$\delta_t = \frac{\partial \ell_t}{\partial h_t} + \left( \frac{\partial h_{t+1}}{\partial h_t} \right)^\top \delta_{t+1}.$$

$$\frac{\partial h_{t+1}}{\partial h_t} = \text{diag}(\phi'(a_{t+1})) W_{hh},$$

# Backpropagation Through Time (BPTT) — derivation

Loss over sequence

$$\mathcal{L} = \sum_{t=1}^T \ell(y_t, \hat{y}_t)$$

We want  $\frac{\partial \mathcal{L}}{\partial W_{hh}}$ . Use chain rule across time:

$$\frac{\partial \mathcal{L}}{\partial W_{hh}} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial h_t} \frac{\partial h_t}{\partial W_{hh}}.$$

$$\delta_t = \frac{\partial \ell_t}{\partial h_t} + \left( \frac{\partial h_{t+1}}{\partial h_t} \right)^\top \delta_{t+1}.$$

$$\frac{\partial h_{t+1}}{\partial h_t} = \text{diag}(\phi'(a_{t+1})) W_{hh},$$

# Backpropagation Through Time (BPTT) — derivation

Loss over sequence

$$\mathcal{L} = \sum_{t=1}^T \ell(y_t, \hat{y}_t)$$

We want  $\frac{\partial \mathcal{L}}{\partial W_{hh}}$ . Use chain rule across time:

$$\frac{\partial \mathcal{L}}{\partial W_{hh}} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial h_t} \frac{\partial h_t}{\partial W_{hh}}.$$

Compute recursive gradient for hidden states. Define  $\delta_t \equiv \frac{\partial \mathcal{L}}{\partial h_t}$ . Then

$$\delta_t = \frac{\partial \ell_t}{\partial h_t} + \left( \frac{\partial h_{t+1}}{\partial h_t} \right)^\top \delta_{t+1}.$$

$$\frac{\partial h_{t+1}}{\partial h_t} = \text{diag}(\phi'(a_{t+1}))W_{hh},$$

# Backpropagation Through Time (BPTT) — derivation

Loss over sequence

$$\mathcal{L} = \sum_{t=1}^T \ell(y_t, \hat{y}_t)$$

We want  $\frac{\partial \mathcal{L}}{\partial W_{hh}}$ . Use chain rule across time:

$$\frac{\partial \mathcal{L}}{\partial W_{hh}} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial h_t} \frac{\partial h_t}{\partial W_{hh}}.$$

Compute recursive gradient for hidden states. Define  $\delta_t \equiv \frac{\partial \mathcal{L}}{\partial h_t}$ . Then

$$\delta_t = \frac{\partial \ell_t}{\partial h_t} + \left( \frac{\partial h_{t+1}}{\partial h_t} \right)^\top \delta_{t+1}.$$

Since  $h_{t+1} = \phi(W_{xh}x_{t+1} + W_{hh}h_t + b)$ ,

$$\frac{\partial h_{t+1}}{\partial h_t} = \text{diag}(\phi'(a_{t+1}))W_{hh},$$

# Backpropagation Through Time (BPTT) — derivation

Loss over sequence

$$\mathcal{L} = \sum_{t=1}^T \ell(y_t, \hat{y}_t)$$

We want  $\frac{\partial \mathcal{L}}{\partial W_{hh}}$ . Use chain rule across time:

$$\frac{\partial \mathcal{L}}{\partial W_{hh}} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial h_t} \frac{\partial h_t}{\partial W_{hh}}.$$

Compute recursive gradient for hidden states. Define  $\delta_t \equiv \frac{\partial \mathcal{L}}{\partial h_t}$ . Then

$$\delta_t = \frac{\partial \ell_t}{\partial h_t} + \left( \frac{\partial h_{t+1}}{\partial h_t} \right)^\top \delta_{t+1}.$$

Since  $h_{t+1} = \phi(W_{xh}x_{t+1} + W_{hh}h_t + b)$ ,

$$\frac{\partial h_{t+1}}{\partial h_t} = \text{diag}(\phi'(a_{t+1}))W_{hh},$$

# Vanishing / Exploding gradients (intuition)

- Expand recurrence:  $\delta_t$  involves products of  $(W_{hh}^\top \text{diag}(\phi'))^k$ .

# Vanishing / Exploding gradients (intuition)

- Expand recurrence:  $\delta_t$  involves products of  $(W_{hh}^\top \text{diag}(\phi'))^k$ .
- If spectral radius  $\rho(W_{hh}) < 1$  (and  $\|\phi'\|$  bounded) gradients decay exponentially → **vanishing**.

# Vanishing / Exploding gradients (intuition)

- Expand recurrence:  $\delta_t$  involves products of  $(W_{hh}^\top \text{diag}(\phi'))^k$ .
- If spectral radius  $\rho(W_{hh}) < 1$  (and  $\|\phi'\|$  bounded) gradients decay exponentially → **vanishing**.
- If  $\rho(W_{hh}) > 1$  gradients grow exponentially → **exploding**.

# Vanishing / Exploding gradients (intuition)

- Expand recurrence:  $\delta_t$  involves products of  $(W_{hh}^\top \text{diag}(\phi'))^k$ .
- If spectral radius  $\rho(W_{hh}) < 1$  (and  $\|\phi'\|$  bounded) gradients decay exponentially → **vanishing**.
- If  $\rho(W_{hh}) > 1$  gradients grow exponentially → **exploding**.
- Remedies: gradient clipping, orthogonal init, truncated BPTT, gated cells (LSTM/GRU).

# Autoencoder setup (general)

$$z = f_\theta(x), \quad \hat{x} = g_\phi(z)$$

$$\mathcal{L}(\theta, \phi) = \frac{1}{N} \sum_{i=1}^N \ell\left(x^{(i)}, g_\phi(f_\theta(x^{(i)}))\right).$$

# Autoencoder setup (general)

$$z = f_\theta(x), \quad \hat{x} = g_\phi(z)$$

$$\mathcal{L}(\theta, \phi) = \frac{1}{N} \sum_{i=1}^N \ell(x^{(i)}, g_\phi(f_\theta(x^{(i)}))).$$

Common choice  $\ell = \text{MSE}$ :  $\|x - \hat{x}\|_2^2$ .

# Linear autoencoder $\Rightarrow$ PCA (sketch)

Assume linear encoder/decoder without biases:

$$z = W_e x, \quad \hat{x} = W_d z = W_d W_e x.$$

$$\min_{W_d, W_e} \|X - W_d W_e X\|_F^2.$$

# Linear autoencoder $\Rightarrow$ PCA (sketch)

Assume linear encoder/decoder without biases:

$$z = W_e x, \quad \hat{x} = W_d z = W_d W_e x.$$

Minimize Frobenius norm over data matrix  $X \in \mathbb{R}^{d \times N}$ :

$$\min_{W_d, W_e} \|X - W_d W_e X\|_F^2.$$

# Linear autoencoder $\Rightarrow$ PCA (sketch)

Assume linear encoder/decoder without biases:

$$z = W_e x, \quad \hat{x} = W_d z = W_d W_e x.$$

Minimize Frobenius norm over data matrix  $X \in \mathbb{R}^{d \times N}$ :

$$\min_{W_d, W_e} \|X - W_d W_e X\|_F^2.$$

Set  $W = W_d W_e$  with  $\text{rank}(W) \leq m$ . By Eckart–Young: best rank- $m$  approximation to  $X$  (in Frobenius norm) is  $X_m = U_m \Sigma_m V_m^\top$ , the PCA truncation.

# Linear autoencoder $\Rightarrow$ PCA (sketch)

Assume linear encoder/decoder without biases:

$$z = W_e x, \quad \hat{x} = W_d z = W_d W_e x.$$

Minimize Frobenius norm over data matrix  $X \in \mathbb{R}^{d \times N}$ :

$$\min_{W_d, W_e} \|X - W_d W_e X\|_F^2.$$

Set  $W = W_d W_e$  with  $\text{rank}(W) \leq m$ . By Eckart–Young: best rank- $m$  approximation to  $X$  (in Frobenius norm) is  $X_m = U_m \Sigma_m V_m^\top$ , the PCA truncation. Thus the column space of optimal  $W$  is span of top- $m$  principal components. With orthonormal constraints  $W_e = U_m^\top, W_d = U_m$  one recovers PCA projection.

# Nonlinear AE: gradient descent

$$\frac{\partial \mathcal{L}}{\partial \theta} = \frac{1}{N} \sum_i \frac{\partial \ell}{\partial \hat{x}} \frac{\partial \hat{x}}{\partial z} \frac{\partial z}{\partial \theta},$$

# Nonlinear AE: gradient descent

$$\frac{\partial \mathcal{L}}{\partial \theta} = \frac{1}{N} \sum_i \frac{\partial \ell}{\partial \hat{x}} \frac{\partial \hat{x}}{\partial z} \frac{\partial z}{\partial \theta},$$

autodiff applies backprop through encoder and decoder stacks.



# LSTM cell equations (forward)

$$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$$

$$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$$

$$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$$

$$\tilde{c}_t = \tanh(W_c x_t + U_c h_{t-1} + b_c)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

$$h_t = o_t \odot \tanh(c_t).$$

# LSTM cell equations (forward)

$$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$$

$$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$$

$$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$$

$$\tilde{c}_t = \tanh(W_c x_t + U_c h_{t-1} + b_c)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

$$h_t = o_t \odot \tanh(c_t).$$

This structure yields controlled gradient flow via multiplicative gates.

# LSTM: gradient flow intuition

- Key term:  $c_t = f_t \odot c_{t-1} + \dots$ . If  $f_t \approx 1$  and  $o_t \approx 1$ , information (and gradients) flow across many steps with little attenuation.


# LSTM: gradient flow intuition

- Key term:  $c_t = f_t \odot c_{t-1} + \dots$ . If  $f_t \approx 1$  and  $o_t \approx 1$ , information (and gradients) flow across many steps with little attenuation.
- Gradient w.r.t.  $c_{t-1}$ :

$$\frac{\partial \mathcal{L}}{\partial c_{t-1}} = \frac{\partial \mathcal{L}}{\partial c_t} \odot f_t + \dots$$

so  $f_t$  acts multiplicatively on gradients, allowing the cell to preserve or forget information dynamically.

# LSTM cell diagram (animated reveal)



# CNN-AE: architecture idea

- Encoder: sequence of Conv + nonlinearity + Pool (downsample)  $\Rightarrow$  latent.

# CNN-AE: architecture idea


- Encoder: sequence of Conv + nonlinearity + Pool (downsample)  $\Rightarrow$  latent.
- Decoder: sequence of Upsample (or transpose-conv) + Conv  $\Rightarrow$  reconstruction.

# Convolution arithmetic (quick notes)

$$\text{output size} = \left\lfloor \frac{H + 2P - K}{S} \right\rfloor + 1,$$

where  $K$  kernel,  $S$  stride,  $P$  padding. Use transpose-convolution (deconv) to invert downsampling or use upsampling + conv.

# CNN Autoencoder diagram



# VAE: goal and generative model

Assume generative model  $p_{\theta}(x, z) = p_{\theta}(x|z)p(z)$  with prior  $p(z) = \mathcal{N}(0, I)$ . The marginal likelihood is:

$$p_{\theta}(x) = \int p_{\theta}(x|z)p(z) dz,$$

intractable in general.

# ELBO derivation (stepwise)

$$\begin{aligned}\log p_\theta(x) &= \log \int p_\theta(x|z)p(z) dz \\&= \log \int q_\phi(z|x) \frac{p_\theta(x|z)p(z)}{q_\phi(z|x)} dz \\&\geq \int q_\phi(z|x) \log \frac{p_\theta(x|z)p(z)}{q_\phi(z|x)} dz \quad (\text{Jensen}) \\&= \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x) \| p(z)) \\&\equiv \mathcal{L}(\theta, \phi; x) \quad (\text{ELBO}).\end{aligned}$$

# ELBO derivation (stepwise)

$$\begin{aligned}\log p_\theta(x) &= \log \int p_\theta(x|z)p(z) dz \\&= \log \int q_\phi(z|x) \frac{p_\theta(x|z)p(z)}{q_\phi(z|x)} dz \\&\geq \int q_\phi(z|x) \log \frac{p_\theta(x|z)p(z)}{q_\phi(z|x)} dz \quad (\text{Jensen}) \\&= \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x) \| p(z)) \\&\equiv \mathcal{L}(\theta, \phi; x) \quad (\text{ELBO}).\end{aligned}$$

Maximizing ELBO  $\approx$  maximizing a lower bound on  $\log p_\theta(x)$ .

# Reparameterization trick

$$z \sim q_\phi(z|x) = \mathcal{N}(\mu_\phi(x), \Sigma_\phi(x))$$

Instead sample  $\epsilon \sim \mathcal{N}(0, I)$  and set

$$z = \mu_\phi(x) + L_\phi(x) \epsilon,$$

where  $L_\phi$  is e.g. lower-triangular matrix s.t.  $L_\phi L_\phi^\top = \Sigma_\phi$ . For diagonal covariances:

$$z = \mu_\phi(x) + \sigma_\phi(x) \odot \epsilon.$$

This makes gradients w.r.t.  $\phi$  tractable via backprop.

# VAE training objective (practical)

For Gaussian decoder (and diagonal posterior) the ELBO per datapoint:

$$\mathcal{L}(\theta, \phi; x) = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - \frac{1}{2} \sum_j (1 + \log \sigma_j^2 - \mu_j^2 - \sigma_j^2).$$

Implement with Monte Carlo estimate (often single sample) and backprop through reparameterization.

# Variational Autoencoder (VAE)

